

BRO CHEAT SHEET

Version: November 17, 2011
Website: <http://www.bro-ids.org>
Email: info@bro-ids.org
Download: <https://github.com/broids/cheat-sheet>



Startup

`bro [options] [file ...]`
`file`Bro policy script or stdin
`-e code`Augment policies by given code
`-h` Display command line options
`-i iface` Read from given interface
`-p pfx` Add given prefix to policy resolution
`-r file` Read from given PCAP file
`-w file` Write to given file in PCAP format
`-x file` Print contents of state file
`-C` Ignore invalid checksum

Language

Lowercase letters represent instance variables and uppercase letters represent types. In general, `x` is an instance of type `T` and `y` an instance of type `U`. Argument names and record fields begin with `a`, `b`, ..., and `z` represents a default instance variable which takes on the type of the right-hand side expression. For notational convenience, `x` can often be replaced with an expression of type `T`.

Variables

Constant qualifier `const`
Constant redefinition `redef x op expr`
Scope qualifier `local, global`
Declaration `scope x: T`
Declaration & Definition `scope z = expr`

Declarations

Type `type name: T`
Function `function f(a: T, ...): R`
Event `event e(a: T, ...)`

Modules

Script import `@load path`
Set current namespace to `ns` `module ns`
Export global symbols `export { ... }`
Access module or enum namespace `T::a`

Statements

Basic statement `stmt`; or `expr`;
Code block `{ stmt; ... }`
Assignment `z = expr`
Function assignment `z = function(...): R { ... }`
Event queuing `event e(...)`
Event scheduling `schedule 10 secs { e(...) }`
Print expression to stdout `print expr`

BRANCHING	ITERATION	CONTROL
<code>if (expr)</code>	<code>for (i in x)</code>	<code>break</code>
<code>{ ... }</code>	<code>{ ... }</code>	<code>continue</code>
<code>else if (expr)</code>	ASYNCHRONOUS	<code>next</code>
<code>{ ... }</code>		<code>return</code>
<code>else</code>	<code>when (expr) { ... }</code>	
<code>{ ... }</code>	<code>when (local x = expr) { ... }</code>	

Expressions

OPERATORS
`!` Negation
`$, ?$` Dereference, record field existence
`+, -, *, /, %` Arithmetic
`++, --` Post-increment, post-decrement
`+=, -=, *=, /=` Arithmetic and assignment
`==, !=` Equality, inequality
`<, <=, >=, >` Less/greater than (or equal)
`&&, ||` Conjunction, disjunction
`in, !in` Membership or pattern matching
`[x]` Index strings and containers
`|x|` Cardinality/size for strings and containers
`f(...)` Function call
`expr ? expr : expr` Ternary if-then-else

Types

BASIC
`addr` IP address (127.0.0.1)

`bool` Boolean flag (T, F)
`count` 64-bit unsigned integer (42)
`double` Double-precision floating point (99.9)
`int` 64-bit signed integer (-7)
`interval` Time interval (8 sec/min/hr/day[s])
`pattern` Regular expression (`/~br[o0])$/`)
`port` Transport-layer port (22/tcp, 53/udp)
`string` String of bytes ("foo")
`subnet` CIDR subnet mask (10.0.0.0/8)
`time` Absolute epoch time (1320977325)

ENUMERABLES

Declaration `enum { FOO, BAR }`
Assignment `scope x = FOO`

RECORDS

Declaration `record { a: T, b: U, ... }`
Constructor `record($a=x, $b=y, ...)`
Assignment `scope r = [$a=x, $b=y, ...]`
Access `z = r$a`
Field assignment `r$b = y`
Deletion `delete r$a`

SETS

Declaration `set[T]`
Constructor `set(x, ...)`
Assignment `scope s = { x, ... }`
Access `z = s[x]`
Insertion `add s[x]`
Deletion `delete s[x]`

TABLES

Declaration `table[T] of U`
Constructor `table([x] = y, ...)`
Assignment `scope t = { [x] = y, ... }`
Access `z = t[x]`
Insertion `t[x] = y`
Deletion `delete t[x]`

VECTORS

Declaration `vector of T`
Constructor `vector(x, ...)`
Assignment `scope v = { x, ... }`
Access `z = v[0]`
Insertion `v[42] = x`

Attributes

Attributes occur at the end of type/event declarations and change their behavior. The syntax is `&key` or `&key=val`, e.g., type `T`: `set[count] &read_expire=5min` or event `foo()` `&priority=-3`.

`&optional` Allow record field to be missing
`&default=x` Use default value `x` for record fields and container elements
`&redef` Allow for redefinition of initial object value
`&expire_func=f` Call `f` right before container element expires
`&read_expire=x` Remove element after not reading it for time `x`
`&write_expire=x` Remove element after not writing it for time `x`
`&create_expire=x` Remove element after time `x` from insertion
`&persistent` Write state to disk (per default on shutdown)
`&synchronized` Synchronize variable across nodes
`&raw_output` Do not escape non-ASCII characters when writing to a file
`&mergeable` Prefer set union to assignment for synchronized state
`&priority=x` Execution priority of event handler, high to low, $x \in [-10, 10]$
`&group="x"` Events in the same group can be jointly activated/deactivated
`&log` Write record field to log

Built-In Functions (BIFs)

Core

- `length(v: any): count`
Returns the number of elements in the container `v`.
- `same_object(o1: any, o2: any): bool`
Check whether `o1` and `o2` reference the same internal object.
- `clear_table(v: any): any`
Remove all elements from the set or table `v`.

Conversion

- `cat(...): string`
Concatenates all given arguments into a single string.
- `cat_sep(sep: string, def: string, ...): string`
Similar to `cat`, but places `sep` between each given argument. *TODO: what does `def` do?*
- `fmt(...): string`
Produces a formatted string. The first argument is the *format string* and specifies how subsequent arguments are converted for output. It is composed of zero or more directives: ordinary characters (not `%`), which are copied unchanged to the output, and conversion specifications, each of which fetches zero or more sub-

sequent arguments. Conversion specifications begin with `%` and the arguments must properly correspond to the specifier.

After the `%`, the following characters may appear in sequence:

<code>%</code>	Literal <code>%</code>
<code>-</code>	Left-align field
<code>[0-9]+</code>	The field width (< 128)
<code>.</code>	Precision of floating point specifiers <code>[efg]</code> (< 128)
<code>A</code>	ALTERNATIVE_STYLE <i>TODO: means what?</i>
<code>[DTdxsefg]</code>	Format specifier
<code>[DT]</code>	ISO timestamp with microsecond precision
<code>d</code>	Signed/Unsigned integer (using C-style <code>%lld/%llu</code> for <code>int/count</code>)
<code>x</code>	Unsigned hexadecimal (using C-style <code>%llx</code>); addresses/ports are converted to host-byte order
<code>s</code>	Escaped string
<code>[efg]</code>	Double

- `type_name(t: any): string`
Returns the type name of `t`.
- `record_type_to_vector(rt: string): vector of string`
Converts the record type name `rt` into a vector of strings, where each element is the name of a record field. Nested records are flattened.
- `to_int(s: string): int`
Converts a `string` into a (signed) integer.
- `int_to_count(n: int): count`
Converts a positive integer into a `count` or returns 0 if $n < 0$.
- `double_to_count(d: double): count`
Converts a positive double into a `count` or returns 0 if $d < 0.0$.
- `to_count(s: string): count`
Converts a `string` into a `count`.
- `interval_to_double(i: interval): double`
Converts an `interval` time span into a `double`.
- `double_to_interval(d: double): interval`
Converts a `double` into an `interval`.
- `time_to_double(t: time): double`
Converts a `time` value into a `double`.
- `double_to_time(d: double): time`
Converts a `double` into a `time` value.
- `double_to_time(d: double): time`
Converts a `double` into a `time` value.
- `port_to_count(p: port): count`
Returns the port number of `p` as `count`.
- `count_to_port(c: count, t: transport_proto): port`
Create a `port` with number `c` and transport protocol `t`.

- `to_port(c count, t: transport_proto): port`
Same as `count_to_port`.
- `addr_to_count(a: addr): count`
Converts an IP address into a 32-bit unsigned integer.
- `count_to_v4_addr(ip: count): addr`
Converts an unsigned integer into an IP address.
- `to_addr(ip: string): addr`
Converts a `string` into an IP address.
- `raw_bytes_to_v4_addr(b: string): addr`
Converts a `string` of bytes into an IP address. It interprets the first 4 bytes of `b` as an IPv4 address in network order.

Math

- `floor(d: double): double`
Chops off any decimal digits, i.e., computes $\lfloor d \rfloor$.

String Processing

- `byte_len(s: string): count`
Returns the number of characters (i.e., bytes) of `s`.
- `sub_bytes(s: string, start: count, n: int): string`
Get a substring of `s`, starting at position `start` and having length `n`.
- `split(s: string, re: pattern): table[count] of string`
Split `s` into an array using `re` to separate the elements. The returned table starts at index 1. Note that conceptually the return value is meant to be a vector and this might change in the future.
- `split1(s: string, re: pattern): table[count] of string`
Same as `split`, but `s` is only split once (if possible) at the earliest position and an array of two strings is returned. An array of one string is returned when `s` cannot be split.
- `split_all(s: string, re: pattern): table[count] of string`
Same as `split`, but also include the matching separators, e.g., `split_all("a-b--cd", /(\-)+/)` returns `{"a", "-", "b", "--", "cd"}`. Odd-indexed elements do not match the pattern and even-indexed ones do.
- `split_n(s: string, re: pattern, incl_sep: bool, max_num_sep: count): table[count] of string`
Similar to `split1` and `split_all`, but `incl_sep` indicates whether to include matching separators and `max_num_sep` the number of times to split `s`.
- `sub(s: string, re: pattern, repl: string): string`
Substitutes `repl` for the first occurrence of `re` in `s`.
- `gsub(s: string, re: pattern, repl: string): string`
Same as `sub` except that all occurrences of `re` are replaced.
- `strcmp(s1: string, s2: string): int` Lexicographically compare `s1` and `s2`. Returns an integer greater than, equal to, or less than 0 according as `s1` is greater than, equal to, or less than `s2`.
- `strstr(big: string, little: string): count`
Locate the first occurrence of `little` in `big`. Returns 0 if `little` is not found in `big`.
- `subst_string(s: string, from: string, to: string): string`
Substitute each (non-overlapping) appearance of `from` in `s` to `to`, and return the resulting string.
- `to_lower(s: string): string`
Returns a copy of `s` with each letter converted to lower case.
- `to_upper(s: string): string`
Returns a copy of `s` with each letter converted to upper case.
- `clean(s: string): string`
Replace non-printable characters in `s` with escaped sequences, with the mappings `0` \rightarrow `\0`, `DEL` \rightarrow `^?`, values ≤ 26 \rightarrow `^[A-Z]`, and values not in `[32,126]` \rightarrow `%XX`.
- `to_string_literal(s: string): string`
Same as `clean`, but with different mappings: values not in `[32,126]` \rightarrow `%XX`, `\` \rightarrow `\\`, `'` \rightarrow `\'`, `"` \rightarrow `\"`.
- `is_ascii(s: string): bool`
Returns false if any byte value of `s` is greater than 127, and true otherwise.
- `escape_string(s: string): string`
Returns a printable version of `s`. Same as `clean` except that non-printable characters are removed.
- `string_to_ascii_hex(s: string): string`
Returns an ASCII hexadecimal representation of a string.
- `str_split(s: string, idx: vector of count): vector of string`
Splits `s` into substrings, taking all the indices in `idx` as cutting points; `idx` does not need to be sorted and out-of-bounds indices are ignored.
- `strip(s: string): string`
Strips whitespace at both ends of `s`.
- `string_fill(len: int, source: string): string`
Generates a string of size `len` and fills it with repetitions of `source`.
- `str_shell_escape(source: string): string`
Takes a string and escapes characters that would allow execution of commands at the shell level. Must be used before including strings in `system()` or similar calls.
- `find_all(s: string, re: pattern) : set of string`
Returns all occurrences of `re` in `s` (or an empty empty set if none).
- `find_last(s: string, re: pattern) : string`
Returns the last occurrence of `re` in `s`. If not found, returns an empty string.

Note that this function returns the match that starts at the largest index in the string, which is not necessarily the longest match. For example, a pattern of `/*` will return the final character in the string.

- **`hexdump(data: string) : string`**

Returns a hex dump for **`data`**. The hex dump renders 16 bytes per line, with hex on the left and ASCII (where printable) on the right. Based on Netdude's hex editor code.